

import the libraries

 Generate


Close

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

read dataset

 Generate



Close

```
btc = pd.read_csv('BTC-USD.csv')
```

15 Basic functions

Double-click (or enter) to edit


```
# 1
print("1. First 5 rows:")
print(btc.head())
```

 1. First 5 rows:

	Date	Open	High	Low	Close	Adj Close	\
0	2014-09-17	465.864014	468.174011	452.421997	457.334015	457.334015	
1	2014-09-18	456.859985	456.859985	413.104004	424.440002	424.440002	
2	2014-09-19	424.102997	427.834991	384.532013	394.795990	394.795990	
3	2014-09-20	394.673004	423.295990	389.882996	408.903992	408.903992	
4	2014-09-21	408.084991	412.425995	393.181000	398.821014	398.821014	

	Volume
0	21056800
1	34483200
2	37919700
3	36863600
4	26580100


```
# 2
print("\n2. Last 5 rows:")
print(btc.tail())
```

 2. Last 5 rows:

	Date	Open	High	Low	Close	\
2708	2022-02-15	42586.464844	44667.218750	42491.035156	44575.203125	
2709	2022-02-16	44578.277344	44578.277344	43456.691406	43961.859375	
2710	2022-02-17	43937.070313	44132.972656	40249.371094	40538.011719	
2711	2022-02-18	40552.132813	40929.152344	39637.617188	40030.976563	
2712	2022-02-19	40022.132813	40246.027344	40010.867188	40126.429688	

	Adj Close	Volume
2708	44575.203125	22721659051
2709	43961.859375	19792547657
2710	40538.011719	26246662813
2711	40030.976563	23310007704
2712	40126.429688	22263900160

```
# 3
print("\n3. Dataset shape:")
print(btc.shape)
```

 3. Dataset shape:

```
(2713, 7)
```

```
# 4
print("\n4. Column names:")
print(btc.columns)
```



```
4. Column names:
Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume'], dtype='object')
```

```
# 5
print("\n5. Data types:")
print(btc.dtypes)
```



```
5. Data types:
Date          object
Open          float64
High          float64
Low           float64
Close         float64
Adj Close     float64
Volume        int64
dtype: object
```

```
# 6
print("\n6. Summary statistics:")
print(btc.describe())
```



```
6. Summary statistics:
      Open      High      Low      Close      Adj Close  \
count  2713.000000  2713.000000  2713.000000  2713.000000  2713.000000
mean   11311.041069  11614.292482  10975.555057  11323.914637  11323.914637
std    16106.428891  16537.390649  15608.572560  16110.365010  16110.365010
min     176.897003   211.731003   171.509995   178.102997   178.102997
25%     606.396973   609.260986   604.109985   606.718994   606.718994
50%     6301.569824  6434.617676   6214.220215   6317.609863   6317.609863
75%    10452.399414  10762.644531  10202.387695  10462.259766  10462.259766
max     67549.734375  68789.625000  66382.062500  67566.828125  67566.828125

      Volume
count  2.713000e+03
mean   1.470462e+10
std    2.001627e+10
min     5.914570e+06
25%     7.991080e+07
50%     5.098183e+09
75%     2.456992e+10
max     3.509679e+11
```

```
# 7
print("\n7. Missing values:")
print(btc.isnull().sum())
```



```
7. Missing values:
Date      0
Open      0
High      0
Low       0
Close     0
Adj Close 0
Volume    0
dtype: int64
```

```
# 8
print("\n8. Unique dates:")
print(btc['Date'].nunique())
```



```
8. Unique dates:
2713
```

```
# 9
print("\n9. Memory usage:")
print(btc.memory_usage())
```



```
9. Memory usage:
Index      132
Date       21704
Open       21704
High       21704
Low        21704
Close      21704
```

```
Adj Close    21704
Volume       21704
dtype: int64
```

```
# 10
print("\n10. Correlation matrix:")
# Select only numeric columns before calculating correlation
print(btc.select_dtypes(include=np.number).corr())
```



```
10. Correlation matrix:
      Open      High      Low      Close  Adj Close  Volume
Open    1.000000  0.999535  0.999103  0.998839  0.998839  0.728537
High    0.999535  1.000000  0.999046  0.999489  0.999489  0.732137
Low      0.999103  0.999046  1.000000  0.999399  0.999399  0.720922
Close    0.998839  0.999489  0.999399  1.000000  1.000000  0.727443
Adj Close 0.998839  0.999489  0.999399  1.000000  1.000000  0.727443
Volume    0.728537  0.732137  0.720922  0.727443  0.727443  1.000000
```

```
# 11
print("\n11. Number of unique values per column:")
print(btc.nunique())
```



```
11. Number of unique values per column:
Date      2713
Open      2709
High      2710
Low       2712
Close     2710
Adj Close 2710
Volume    2713
dtype: int64
```

```
# 12
print("\n12. Index information:")
print(btc.index)
```



```
12. Index information:
RangeIndex(start=0, stop=2713, step=1)
```

```
# 13
print("\n13. Sample of 3 random rows:")
print(btc.sample(3))
```



```
13. Sample of 3 random rows:
      Date      Open      High      Low      Close \
2526  2021-08-17  45936.457031  47139.570313  44512.417969  44695.359375
2210  2020-10-05  10676.529297  10793.507813  10634.600586  10793.339844
1413  2018-07-31   8181.200195   8181.529785   7696.930176   7780.439941

      Adj Close      Volume
2526  44695.359375  33451362600
2210  10793.339844  47537578009
1413   7780.439941   5287530000
```

```
# 14
print("\n14. Count of non-NA values:")
print(btc.count())
```



```
14. Count of non-NA values:
Date      2713
Open      2713
High      2713
Low       2713
Close     2713
Adj Close 2713
Volume    2713
dtype: int64
```




[Close](#)

```
# 15
print("\n15. Value counts for Volume (top 5):")
print(btc['Volume'].value_counts().head())
```



```
15. Value counts for Volume (top 5):
Volume
```

```

22263900160    1
21056800      1
34483200      1
37919700      1
36863600      1
Name: count, dtype: int64

```

✓ Data Cleaning

```

# 1
# Check for missing values
print("Missing values before cleaning:")
print(btc.isnull().sum())

```

```

↔ Missing values before cleaning:
Date          0
Open          0
High          0
Low           0
Close         0
Adj Close     0
Volume        0
dtype: int64

```

```

# Drop rows with any missing values
btc_clean = btc.dropna()

# Convert Date to datetime format
btc_clean['Date'] = pd.to_datetime(btc_clean['Date'])

```

```

# 2
# Verify data types
print("\nData types after cleaning:")
print(btc_clean.dtypes)

```

```

↔ Data types after cleaning:
Date          datetime64[ns]
Open          float64
High          float64
Low           float64
Close         float64
Adj Close     float64
Volume        int64
dtype: object

```

Start coding or [generate](#) with AI.

```

# 3
# Check for duplicates
print("\nNumber of duplicate rows:", btc_clean.duplicated().sum())

```

```

↔ Number of duplicate rows: 0

```

```

# 4
# Remove potential outliers using z-score (for Close price)
z_scores = np.abs(stats.zscore(btc_clean['Close']))
btc_clean = btc_clean[(z_scores < 3)]

print("\nDataset shape after cleaning:", btc_clean.shape)

```

```

↔ Dataset shape after cleaning: (2609, 7)

```

✓ Data Filtering

```

#1
# Filter for prices above $1000
high_prices = btc_clean[btc_clean['Close'] > 1000]
print("Number of days with price > $1000:", len(high_prices))

```

```

↔ Number of days with price > $1000: 1736

```

```
# 2
# Filter for specific year (2017)
btc_2017 = btc_clean[btc_clean['Date'].dt.year == 2017]
print("\n2017 data shape:", btc_2017.shape)
```



2017 data shape: (365, 7)

✓ Grouping

```
# 1
# Group by year and get average closing price
yearly_avg = btc_clean.groupby(btc_clean['Date'].dt.year)['Close'].mean()
print("Yearly average closing prices:")
print(yearly_avg)
```



Yearly average closing prices:

Date	
2014	363.693085
2015	272.453381
2016	568.492407
2017	4006.033629
2018	7572.298947
2019	7395.246282
2020	11116.378092
2021	42661.042460
2022	41345.687735

Name: Close, dtype: float64

```
# 2
# Group by month in 2017 and get max volume
btc_2017 = btc_clean[btc_clean['Date'].dt.year == 2017]
monthly_max_vol = btc_2017.groupby(btc_2017['Date'].dt.month)['Volume'].max()
print("\nMonthly max volume in 2017:")
print(monthly_max_vol)
```



Monthly max volume in 2017:

Date	
1	510199008
2	407220000
3	706598976
4	580444032
5	2406700032
6	2569530112
7	2249260032
8	3764239872
9	4148069888
10	3615480064
11	11568799744
12	22197999616

Name: Volume, dtype: int64

✓ Sorting

```
# 1
# Sort by highest closing price
highest_prices = btc_clean.sort_values('Close', ascending=False)
print("Top 5 highest closing prices:")
print(highest_prices[['Date', 'Close']].head())
```



Top 5 highest closing prices:

	Date	Close
2414	2021-04-27	55033.117188
2579	2021-10-09	54968.222656
2415	2021-04-28	54824.703125
2365	2021-03-09	54824.117188
2628	2021-11-27	54815.078125

```
# 2
# Sort by date (chronological order)
btc_sorted = btc_clean.sort_values('Date')
print("\nFirst 5 dates in chronological order:")
print(btc_sorted[['Date', 'Close']].head())
```



First 5 dates in chronological order:

Date	Close
------	-------

```
0 2014-09-17 457.334015
1 2014-09-18 424.440002
2 2014-09-19 394.795990
3 2014-09-20 408.903992
4 2014-09-21 398.821014
```

✓ Aggregations

```
# Calculate various aggregations for Close price
close_stats = {
    'Minimum': btc_clean['Close'].min(),
    'Maximum': btc_clean['Close'].max(),
    'Mean': btc_clean['Close'].mean(),
    'Median': btc_clean['Close'].median(),
    'Mode': btc_clean['Close'].mode()[0],
    'Variance': btc_clean['Close'].var(),
    'Std Dev': btc_clean['Close'].std()
}

print("Close Price Statistics:")
for stat, value in close_stats.items():
    print(f"{stat}: {value:.2f}")
```

```
↔ Close Price Statistics:
Minimum: 178.10
Maximum: 55033.12
Mean: 9406.60
Median: 5725.59
Mode: 236.15
Variance: 173628838.32
Std Dev: 13176.83
```

✓ Additional Queries

```
# 1. Days with highest price volatility (difference between High and Low)
btc_clean['Volatility'] = btc_clean['High'] - btc_clean['Low']
print("Top 5 most volatile days:")
print(btc_clean.sort_values('Volatility', ascending=False)[['Date', 'Volatility']].head())
```

```
↔ Top 5 most volatile days:
      Date    Volatility
2436 2021-05-19 12864.621094
2635 2021-12-04 11030.062500
2547 2021-09-07  9568.558594
2351 2021-02-23  8914.339844
2429 2021-05-12  8788.828125
```

```
# 2. Percentage change from previous day
btc_clean['Daily_Change'] = btc_clean['Close'].pct_change() * 100
print("\nLargest single day percentage gains:")
print(btc_clean.sort_values('Daily_Change', ascending=False)[['Date', 'Daily_Change']].head())
```

```
↔ Largest single day percentage gains:
      Date    Daily_Change
1177 2017-12-07    25.247169
1037 2017-07-20    23.936087
1176 2017-12-06    19.928334
2336 2021-02-08    18.746474
2010 2020-03-19    18.187756
```

```
# 3. Days where closing price was higher than opening price
up_days = btc_clean[btc_clean['Close'] > btc_clean['Open']]
print("\nNumber of days closing higher than opening:", len(up_days))
```

```
↔ Number of days closing higher than opening: 1408
```

```
# 4. Average volume by year
yearly_volume = btc_clean.groupby(btc_clean['Date'].dt.year)['Volume'].mean()
print("\nAverage yearly volume:")
print(yearly_volume)
```

```
↔ Average yearly volume:
```

```
Date
2014    2.383690e+07
2015    3.390557e+07
2016    8.592451e+07
2017    2.382867e+09
2018    6.063552e+09
2019    1.673049e+10
2020    3.302327e+10
2021    4.676478e+10
2022    2.718684e+10
Name: Volume, dtype: float64
```

```
# 5. Correlation between volume and price change
correlation = btc_clean['Volume'].corr(btc_clean['Daily_Change'])
print("\nCorrelation between volume and daily change:", correlation)
```



```
Correlation between volume and daily change: 0.009254312543180031
```

```
# 6. Days with highest trading volume
print("\nTop 5 days by trading volume:")
print(btc_clean.sort_values('Volume', ascending=False)[['Date', 'Volume']].head())
```



```
Top 5 days by trading volume:
      Date      Volume
2354 2021-02-26  350967941479
2436 2021-05-19  126358098747
2308 2021-01-11  123320567399
2326 2021-01-29  117894572511
2351 2021-02-23  106102492824
```

```
# 7. Monthly average price in 2017
btc_2017 = btc_clean[btc_clean['Date'].dt.year == 2017]
monthly_avg_2017 = btc_2017.groupby(btc_2017['Date'].dt.month)['Close'].mean()
print("\n2017 monthly average prices:")
print(monthly_avg_2017)
```



```
2017 monthly average prices:
Date
1      914.916159
2     1062.533672
3     1129.365228
4     1206.641007
5     1895.383529
6     2636.204346
7     2519.418386
8     3880.989998
9     4064.836312
10    5360.071604
11    7813.132975
12   15294.270980
Name: Close, dtype: float64
```

```
# 8. Price range for each year
yearly_range = btc_clean.groupby(btc_clean['Date'].dt.year)['Close'].agg(['min', 'max'])
print("\nYearly price ranges:")
print(yearly_range)
```



```
Yearly price ranges:
      min      max
Date
2014  310.737000  457.334015
2015  178.102997  465.321014
2016  364.330994  975.921021
2017  777.757019 19497.400391
2018  3236.761719 17527.000000
2019  3399.471680 13016.231445
2020  4970.788086 29001.720703
2021  29374.152344 55033.117188
2022  35030.250000 47686.812500
```

```
# 9. Days where price doubled from previous day
doubled_days = btc_clean[btc_clean['Daily_Change'] > 100]
print("\nDays where price more than doubled:")
print(doubled_days[['Date', 'Daily_Change']])
```



```
Days where price more than doubled:
Empty DataFrame
Columns: [Date, Daily_Change]
Index: []
```

```
# 10. Rolling 7-day average price
btc_clean['7_Day_Avg'] = btc_clean['Close'].rolling(window=7).mean()
print("\n7-day rolling average (last 5 days):")
print(btc_clean[['Date', 'Close', '7_Day_Avg']].tail())
```



```
7-day rolling average (last 5 days):
      Date      Close      7_Day_Avg
2708 2022-02-15  44575.203125  43130.850446
2709 2022-02-16  43961.859375  43077.002232
2710 2022-02-17  40538.011719  42644.559152
2711 2022-02-18  40030.976563  42304.993304
2712 2022-02-19  40126.429688  42002.416295
```

```
# 11. Quarterly performance
btc_clean['Quarter'] = btc_clean['Date'].dt.quarter
quarterly_perf = btc_clean.groupby([btc_clean['Date'].dt.year, 'Quarter'])['Close'].last().unstack()
print("\nQuarterly closing prices:")
print(quarterly_perf)
```



```
Quarterly closing prices:
Quarter      1          2          3          4
Date
2014         NaN         NaN  386.944000  320.192993
2015    244.223999  263.071991  236.059998  430.566986
2016    416.729004  673.336975  609.734985  963.742981
2017    1071.790039  2480.840088  4338.709961  14156.400391
2018    6973.529785  6404.000000  6625.560059  3742.700439
2019    4105.404297  10817.155273  8293.868164  7193.599121
2020    6438.644531  9137.993164  10784.491211  29001.720703
2021    51704.160156  35040.835938  43790.894531  46306.445313
2022    40126.429688         NaN         NaN         NaN
```

```
# 12. Days with lowest volatility
print("\nTop 5 least volatile days:")
print(btc_clean.sort_values('Volatility')[['Date', 'Volatility']].head())
```



```
Top 5 least volatile days:
      Date  Volatility
382 2015-10-04    1.028000
253 2015-05-28    1.172012
262 2015-06-06    1.339996
365 2015-09-17    1.359009
368 2015-09-20    1.455001
```

```
# 13. Price change by year
annual_change = btc_clean.groupby(btc_clean['Date'].dt.year)['Close'].agg(['first', 'last'])
annual_change['Change'] = (annual_change['last'] - annual_change['first']) / annual_change['first'] * 100
print("\nAnnual price changes:")
print(annual_change)
```



```
Annual price changes:
      first      last      Change
Date
2014  457.334015  320.192993  -29.987059
2015  314.248993  430.566986   37.014595
2016  434.334015  963.742981  121.889824
2017   998.325012  14156.400391  1318.015198
2018  13657.200195  3742.700439  -72.595405
2019   3843.520020  7193.599121   87.161745
2020   7200.174316  29001.720703  302.791925
2021  29374.152344  46306.445313   57.643512
2022  47686.812500  40126.429688  -15.854242
```

```
# 14. Volume vs Price scatter
high_volume_days = btc_clean[btc_clean['Volume'] > btc_clean['Volume'].quantile(0.9)]
print("\nAverage price on high volume days:", high_volume_days['Close'].mean())
```



```
Average price on high volume days: 28415.67222864751
```

```
# 15. Price momentum (3-day)
btc_clean['3_Day_Momentum'] = btc_clean['Close'].diff(3)
print("\nRecent 3-day momentum values:")
print(btc_clean[['Date', 'Close', '3_Day_Momentum']].tail())
```




```
Recent 3-day momentum values:
```

	Date	Close	3_Day_Momentum
2708	2022-02-15	44575.203125	2330.734375
2709	2022-02-16	43961.859375	1764.343750
2710	2022-02-17	40538.011719	-2048.906250
2711	2022-02-18	40030.976563	-4544.226562
2712	2022-02-19	40126.429688	-3835.429687

```
# 16. Price bands (20-day)
btc_clean['20_Day_MA'] = btc_clean['Close'].rolling(window=20).mean()
btc_clean['Upper_Band'] = btc_clean['20_Day_MA'] + (btc_clean['Close'].rolling(window=20).std() * 2)
btc_clean['Lower_Band'] = btc_clean['20_Day_MA'] - (btc_clean['Close'].rolling(window=20).std() * 2)
print("\nPrice bands (last 5 days):")
print(btc_clean[['Date', 'Close', '20_Day_MA', 'Upper_Band', 'Lower_Band']].tail())
```



```
Price bands (last 5 days):
```

	Date	Close	20_Day_MA	Upper_Band	Lower_Band
2708	2022-02-15	44575.203125	40877.074414	46372.734271	35381.414557
2709	2022-02-16	43961.859375	41218.255664	46582.262366	35854.248963
2710	2022-02-17	40538.011719	41355.939649	46485.040528	36226.838769
2711	2022-02-18	40030.976563	41450.579492	46396.260739	36504.898246
2712	2022-02-19	40126.429688	41561.020899	46267.374226	36854.667571

```
# 17. Relative Strength Index (14-day)
delta = btc_clean['Close'].diff()
gain = delta.where(delta > 0, 0)
loss = -delta.where(delta < 0, 0)
avg_gain = gain.rolling(window=14).mean()
avg_loss = loss.rolling(window=14).mean()
rs = avg_gain / avg_loss
btc_clean['RSI'] = 100 - (100 / (1 + rs))
print("\nRSI values (last 5 days):")
print(btc_clean[['Date', 'Close', 'RSI']].tail())
```



```
RSI values (last 5 days):
```

	Date	Close	RSI
2708	2022-02-15	44575.203125	71.108013
2709	2022-02-16	43961.859375	77.730362
2710	2022-02-17	40538.011719	60.666638
2711	2022-02-18	40030.976563	43.885888
2712	2022-02-19	40126.429688	44.547513

```
# 18. Days closing above opening by more than 10%
big_up_days = btc_clean[(btc_clean['Close'] - btc_clean['Open']) / btc_clean['Open'] > 0.1]
print("\nDays closing >10% above open:", len(big_up_days))
```



```
Days closing >10% above open: 47
```

```
# 19. Monthly returns
monthly_returns = btc_clean.groupby([btc_clean['Date'].dt.year, btc_clean['Date'].dt.month])['Close'].last().pct_change()
print("\nMonthly returns (last 5 months):")
print(monthly_returns.tail())
```



```
Monthly returns (last 5 months):
```

Date	Date	
2021	10	0.250753
	11	0.000794
	12	-0.155224
2022	1	-0.168947
	2	0.042702

Name: Close, dtype: float64

```
# 20. Volatility clusters
btc_clean['High_Volatility'] = btc_clean['Volatility'] > btc_clean['Volatility'].quantile(0.9)
print("\nHigh volatility days count:", btc_clean['High_Volatility'].sum())
```




```
High volatility days count: 261
```

✓ Data Visualization


```
# Set style
# Use an available seaborn style from matplotlib
```

```
plt.style.use('seaborn-v0_8')
plt.figure(figsize=(12, 6))


# 1. Price Trend Over Time
plt.subplot(2, 3, 1)
plt.plot(btc_clean['Date'], btc_clean['Close'])
plt.title('Bitcoin Price Trend')
plt.xlabel('Date')
plt.ylabel('Price (USD)')
plt.yscale('log') # Log scale for better visualization
```



```
# 2. Daily Price Changes Distribution
plt.subplot(2, 3, 2)
sns.histplot(btc_clean['Daily_Change'].dropna(), bins=50, kde=True)
plt.title('Distribution of Daily Price Changes')
plt.xlabel('Daily Change (%)')
```





```
# 3. Volume vs Price Scatter
plt.subplot(2, 3, 3)
plt.scatter(btc_clean['Volume'], btc_clean['Close'], alpha=0.3)
plt.title('Trading Volume vs Price')
plt.xlabel('Volume')
plt.ylabel('Price (USD)')
plt.yscale('log')
plt.xscale('log')
```



```
# 4. Monthly Average Prices (2017-2018)
btc_recent = btc_clean[btc_clean['Date'].dt.year >= 2017]
monthly_avg = btc_recent.groupby([btc_recent['Date'].dt.year, btc_recent['Date'].dt.month])['Close'].mean().unstack(0)
plt.subplot(2, 3, 4)
monthly_avg.plot(kind='bar')
plt.title('Monthly Average Prices (2017-2018)')
```

```
plt.xlabel('Month')
plt.ylabel('Average Price (USD)')
```

```
Text(0, 0.5, 'Average Price (USD)')
```



```
# 5. Moving Averages
plt.subplot(2, 3, 5)
plt.plot(btc_clean['Date'], btc_clean['Close'], label='Daily Close')
plt.plot(btc_clean['Date'], btc_clean['20_Day_MA'], label='20-Day MA')
plt.plot(btc_clean['Date'], btc_clean['Upper_Band'], label='Upper Band', linestyle='--')
plt.plot(btc_clean['Date'], btc_clean['Lower_Band'], label='Lower Band', linestyle='--')
plt.title('Price with Moving Averages')
plt.xlabel('Date')
plt.ylabel('Price (USD)')
plt.yscale('log')
plt.legend()
```

```
plt.tight_layout()
plt.show()
```

